Semantic tagging

Semantic tagging on Temco hardware · built Nov 13, 2026

Semantic tagging on Temco hardware

T3000 does not have to learn a second name for a point. A point keeps the name it was given, and carries tags beside it: what it measures, what it belongs to, and which equipment it is part of. Those tags are what makes a flat point list readable by a machine — by our own assistant, by Claude or Copilot through MCP, and by a fault-detection rule that has to know which sensor it is looking at.

This is a subject that crosses products: the tags are stored on the controller, served by the web front end, and consumed by anything that speaks MCP.

The order to read it in

1. **How tagging works now** — the current implementation, and the shape of the data.
2. **Auto-tagging** — the rules that fill tags in from a point's own name.
3. **The MCP server** — how an assistant asks for the semantic model and gets it back.
4. **The legacy design** — the single-table version, kept for the controllers still on it.
5. **The assistant and the rules** — the local model, its API, the chat screen, and the FDD engine that reads the same tags.

The plan

The last chapter is the strategic plan this work is being built against: OpenBIM (IFC through Bonsai and thatOpen) together with Haystack, across the C++ front end, the web front end and the firmware.